

The evaluation map — three questions, four datasets, one system

What is evaluated

Which dataset feeds it

Metric & headline number

DEPTH
reconstruction

ClearGrasp val (100 imgs)
corrupted-GT protocol

RMSE (m)
inpaint 0.123 vs solver 0.136

GRASP
targeting

ReSort-IT val (n=251)
native + composite mix

ON-TARGET %
7.2% overall · 11.3 vs 3.1

INSTRUCTION
robustness

ReSort-IT val + held-out
paraphrase / absent / null

parse-fail %
Run2 18→76 · Run3 ~1

REAL-IMAGE
transfer

ZeroWaste val (572 real
MRF conveyor frames)

policy 0% · perception
hit-precision ~30%